~~Billion~~ Trillion-Scale ~~Vector~~ Search on Object Storage

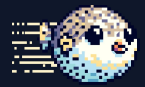

turbopuffer

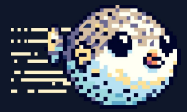

turbopuffer

fast search on object storage @ 10M WPS scale

semantic search
vector similarity

full-text search
traditional search by keywords

aggregations & group by
real database queries

object storage (S3 / GCS), with adaptable SSD/RAM caching
cheap and scalable

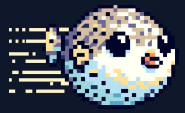

turbopuffer

who’s puffin’

CURSOR

Notion

Linear

Tolan

Pylon

Readwise

top 3 ai lab

SUPERHUMAN

warp

clay

TELUS

.. and many more (white text is annoying, ran out of time)

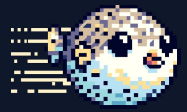

turbopuffer

if you GMI as a new DB, 2 ingredients needed:

1. **New Workload:** Connect LLMs to data
   Vectors are a 10-30x size amplification
   1kb text → 4x 1024d vectors → ~16kb vectors

2. **New Storage Architecture:** Object-Storage Native
   a. **NVMe SSDs in Cloud (2017)**
   b. **S3 Strongly Consistent (2020)**
   c. **S3 CAS (2024)**

(every successful database will support every SQL query eventually)

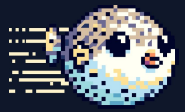

turbopuffer

first object storage native database

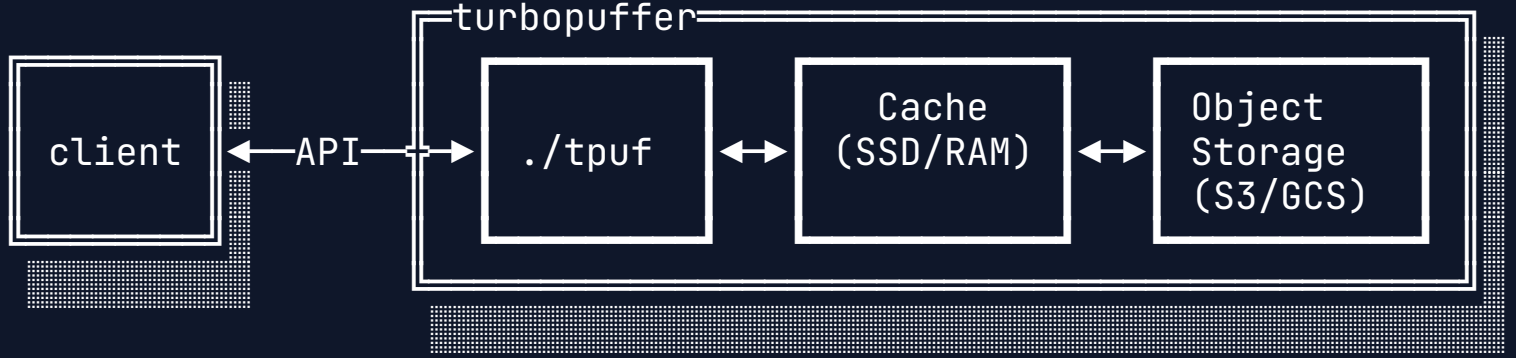

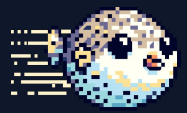

turbopuffer

economics for object storage native databases

Configuration (storage only)	$/GB (USD)
1 × RAM (100% full)	$5
3 × SSD (50% full)	$0.60
2 × SSD (50% full)	$0.40
1 × SSD cache (100%) + S3†	$0.12
S3†	$0.02

† Storage compute separation

Assuming ~$0.10/GB for NVMe and EBS, ~$0.02/GB for S3

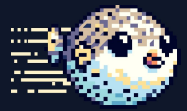

turbopuffer

roundtrip sensitive database w/ high concurrency

Medium	Random Read	Throughput
RAM	100 ns	25 GB/s
NVMe SSD	100 µs	10 GB/s
EBS†	1 ms	5 GB/s
S3 Hedged p99†	200 ms	5 GB/s

† Compete for network bandwidth (on some machines)

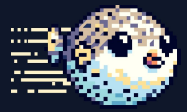

turbopuffer

vector search indexes on object storage

Raw Vectors

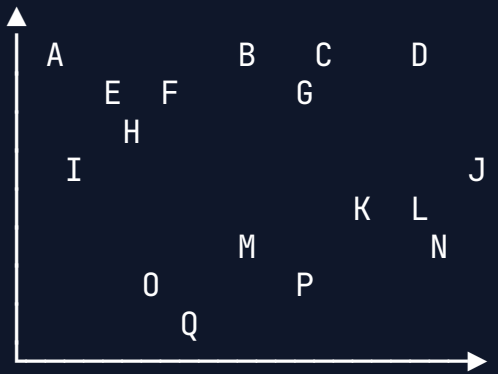

Graph Index

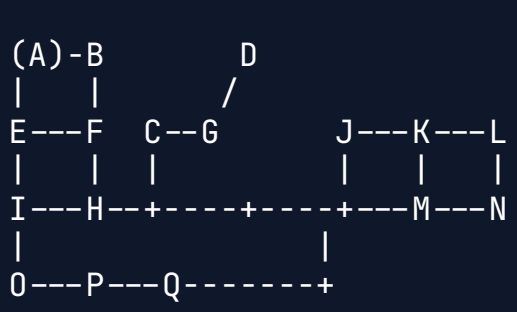

≤ 6 Roundtrips

Clustered Index

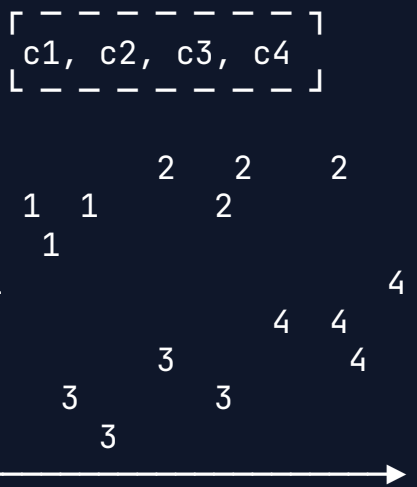

2 Roundtrips

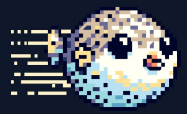

turbopuffer

trade-offs for an object-storage first database

✅ STRENGTHS

- Low cost
- Simple → reliable and horizontally scalable
- Fast warm queries
- High write throughput

⚠ LIMITATIONS

- Cold queries slow
  - ↳ Mitigation: keep full index warm on lower-cost SSDs
- Higher write latency

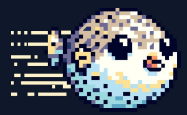

turbopuffer

architecture you want to be oncall for

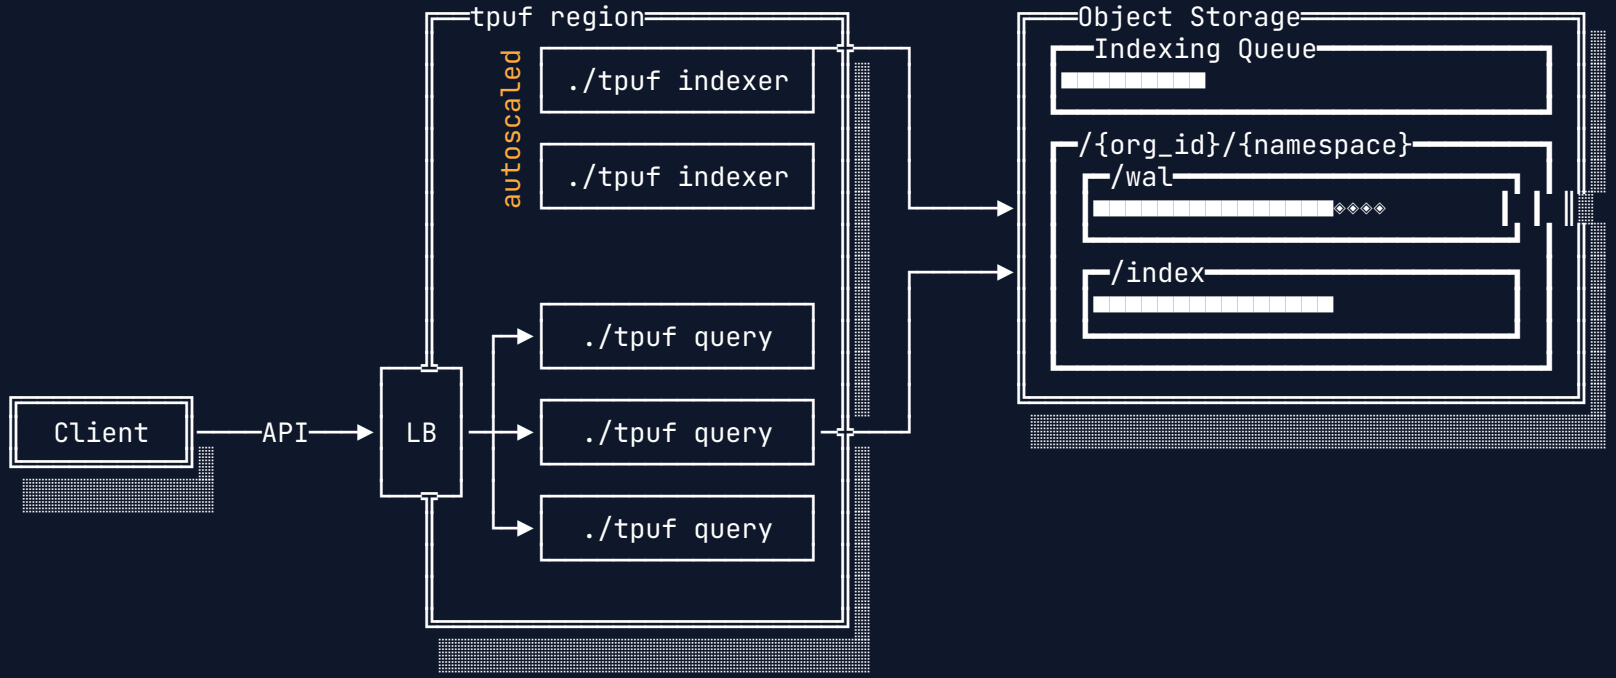

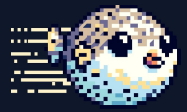

turbopuffer

performance

pctile	Warm namespace	Cold namespace
p50	8 ms	343 ms
p90	10 ms	444 ms
p99	35 ms	554 ms

VECTOR SEARCH – 768 dimensions, 1M docs, ~3GB

pctile	Warm namespace	Cold namespace
p50	11 ms	221 ms
p90	18 ms	285 ms
p99	40 ms	433 ms

FULL-TEXT SEARCH – BM25, 1M docs, ~300MB

turbopuffer

security

all data at rest is encrypted using AES-256

turbopuffer also supports customer managed encryption keys (CMEK)

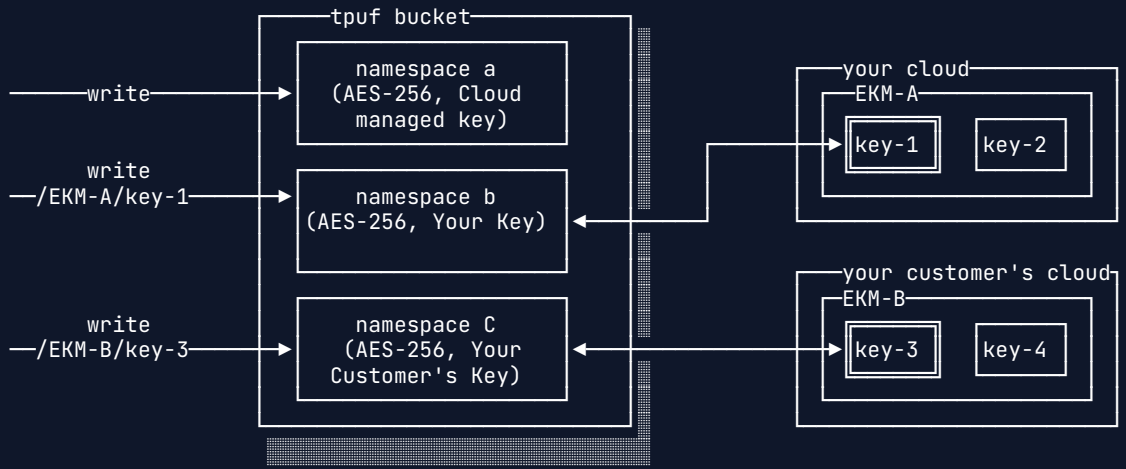

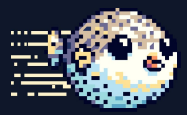

CASE STUDIES

CURSOR

> After switching our vector db to @turbopuffer, we're saving an order of magnitude in costs and dealing with far less complexity!
>
> -Aman Sanger, Co-founder

95%
cost reduction

100B+
vectors

10GB/s
write peaks

10M+
namespaces

* namespace-per-codebase

* active codebase namespaces are loaded into memory/NVMe, inactive fade into object storage

* unlimited namespaces in a fully serverless model; no more bin-packing codebase vector indexes to servers

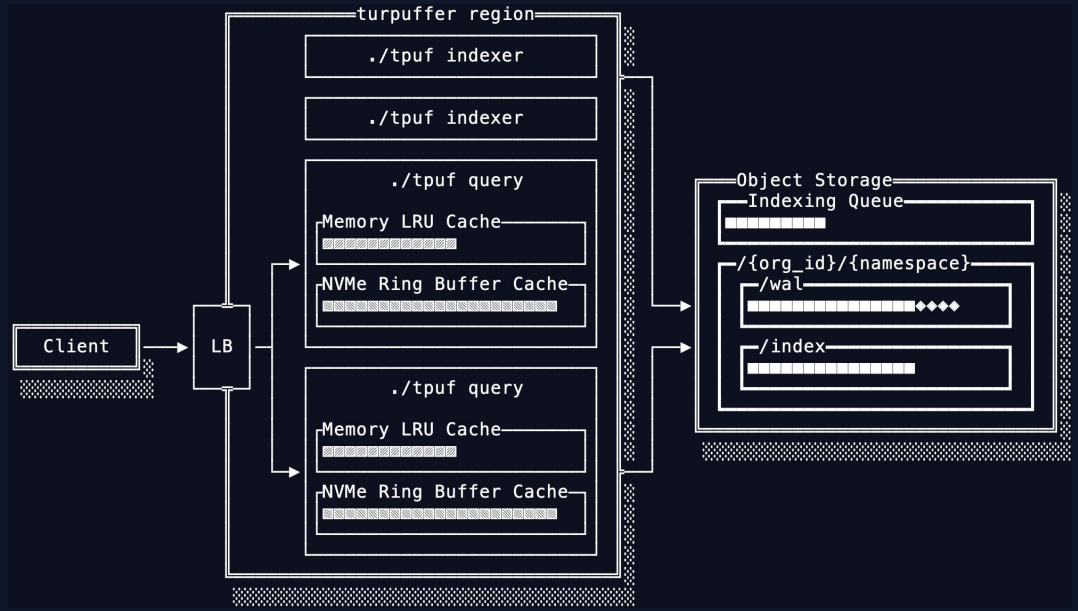

Notion

> turbopuffer's economics have changed the way we think about building products that connect data to users and LLMs.
>
> -Akshay Kothari, Co-founder

millions
$ saved annually

10B+
vectors

1GB/s
write peaks

1M+
namespaces

* Consistent reads with 100,000+ writes/s peaks
* 80% reduction in cost, allowing Notion to remove per-user AI charges
* ⩾ 99.99% uptime
* From concerned to excited about 10x'ing their data size
* Zero performance drops
* A turbopuffer team so responsive they felt part of the Notion engineering team
* A roadmap aligned with their anticipated needs

Linear

> Their responsiveness and shipping velocity make us feel like we are their only customer.
>
> -Tom Moor, Head of Engineering

70%
cost reduction

250M+
documents

13ms
P50 latency

1.5M+
namespaces

* replaced Elasticsearch & pg_vector

* zero-ops search for terabytes of data

* results from parallel queries, using vector + FTS, passed into a reranker

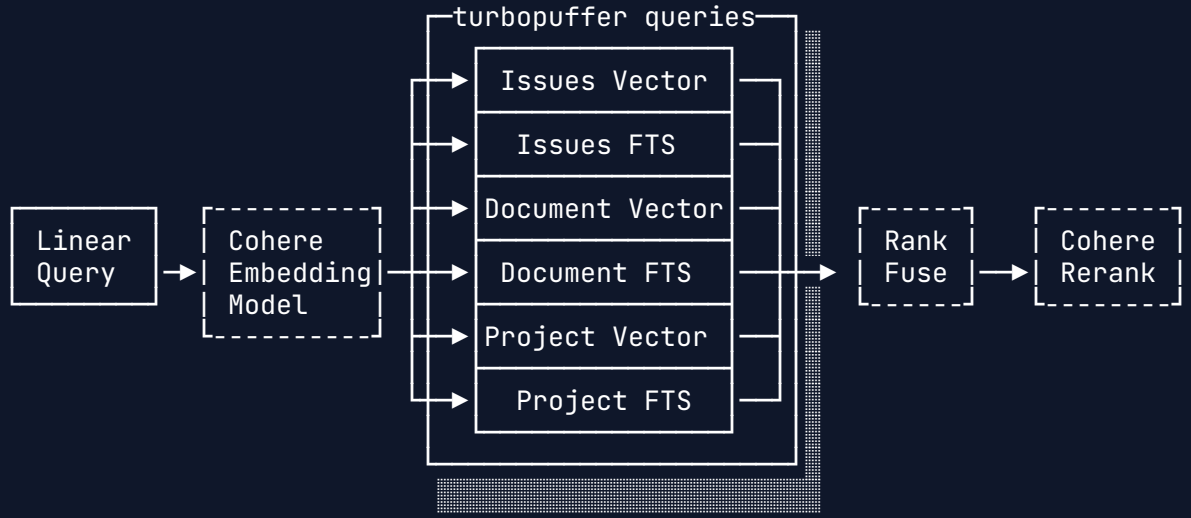

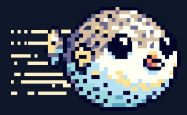

turbopuffer

recall

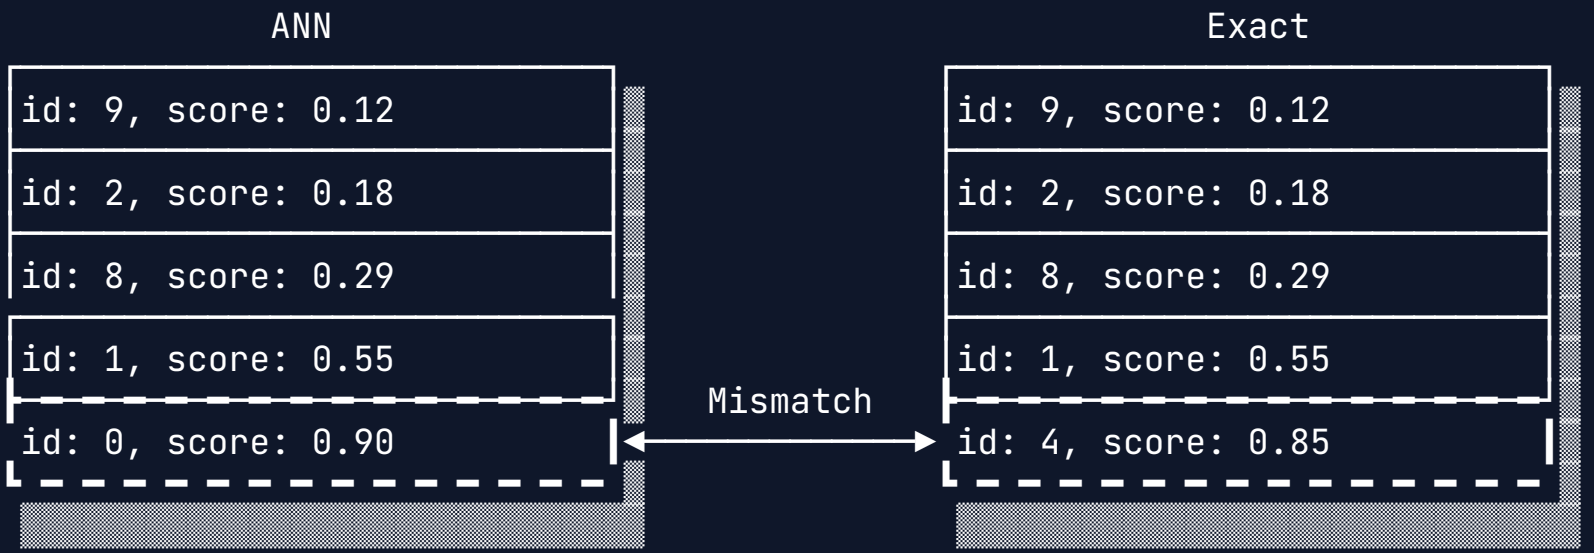

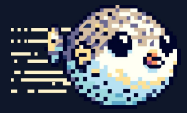

turbopuffer

recall observability

ORG_ID	↑ AVG RECALL@TOPK	AVG RECALL@10/TOPK
[redacted]	97.6 %	98.0 %
[redacted]	98.5 %	90.0 %
[redacted]	99.0 %	99.6 %
[redacted]	99.5 %	100.0 %
[redacted]	99.9 %	100.0 %
[redacted]	100.0 %	100.0 %
[redacted]	100.0 %	100.0 %
[redacted]	100.0 %	100.0 %
[redacted]	100.0 %	100.0 %

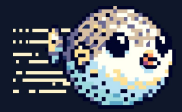

turbopuffer

Q&A

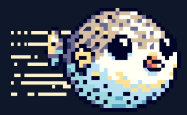

APPENDIX

turbopuffer

filtered recall (prefiltering vs postfiltering vs native)

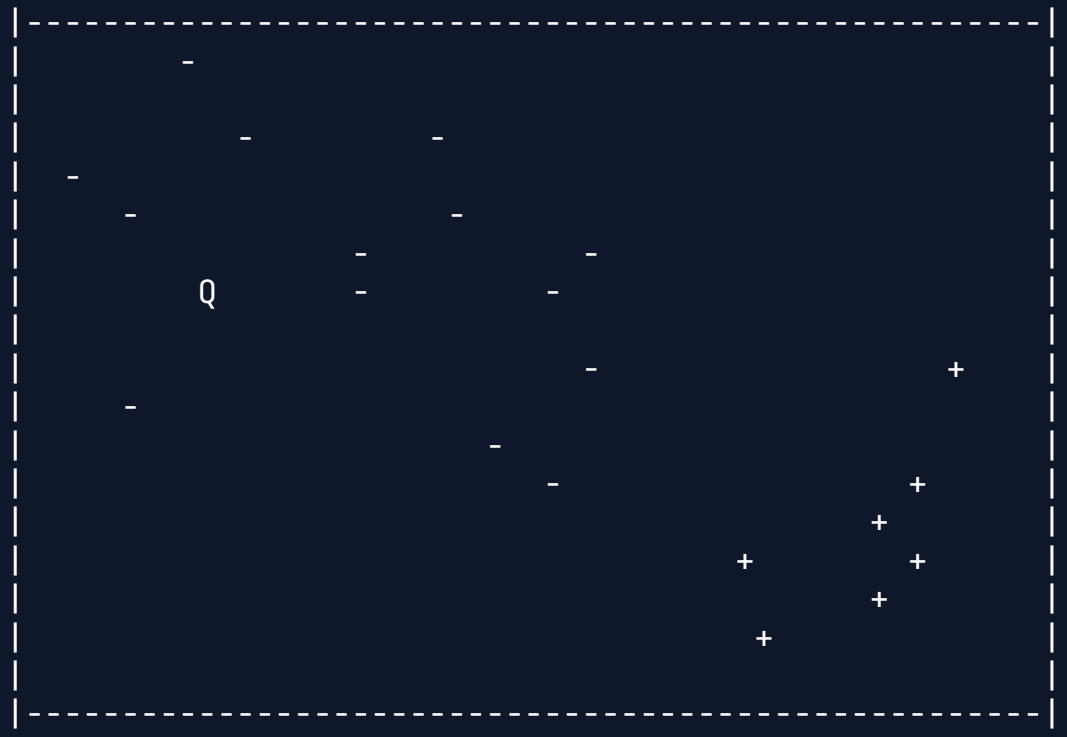

legend:
 Q: query vector
 +: document matching the filter
 -: document not matching the filter